

# Big Transfer (BiT): General Visual Representation Learning

Dmitry Medvedev

Moscow State University

# BiT: General Idea

[Kolesnikov et al., 2020]

**Big Model & Big Data**

**Fine-tuning for visual task of any size**

# Architectures

layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112	7×7, 64, stride 2				
conv2_x	56×56	3×3 max pool, stride 2				
		$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
conv3_x	28×28	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$
conv4_x	14×14	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
conv5_x	7×7	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
	1×1	average pool, 1000-d fc, softmax				
FLOPs		$1.8 \times 10^9$	$3.6 \times 10^9$	$3.8 \times 10^9$	$7.6 \times 10^9$	$11.3 \times 10^9$

# Upstream Pre-Training: Data

Модель	Данные	Объем	Число классов	Label per Image
BiT-L	JFT-300M	300 M	19 K	$\sim 1.26$
BiT-M	ImageNet-21k	14.2 M	21 K	$\geq 1$
BiT-S	ILSVRC-2012	1.28 M	1 K	1


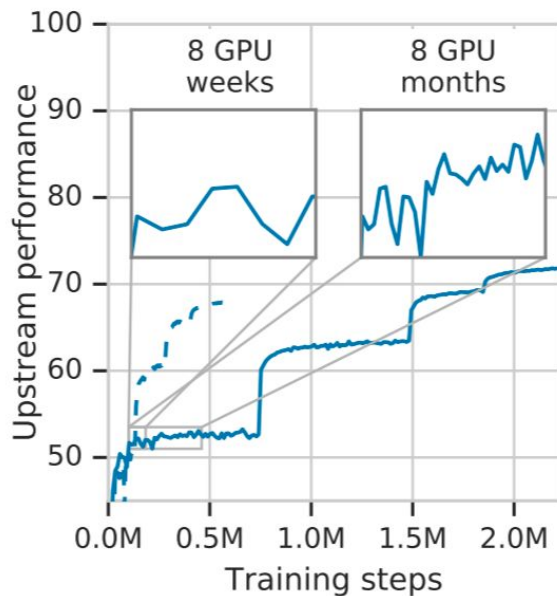

# Upstream Pre-Training

- Scale
- Batch Normalization -> Group Normalization + Weight Standardization
- SGD + momentum
- Аугментации: Crop + Horizontal Flip

Model	Num. Epochs	lr decay by 10 after epochs
-S & -M	90	30, 60, 80
-L	40	23, 30, 37


Image Size	Warm-up steps	Weight decay	Batch Size	Images per chip
224 x 224	5000	1e-4	4096	8

# Upstream Pre-Training





# Batch Normalization

[Santurkar et al. 2020]




# Batch Normalization

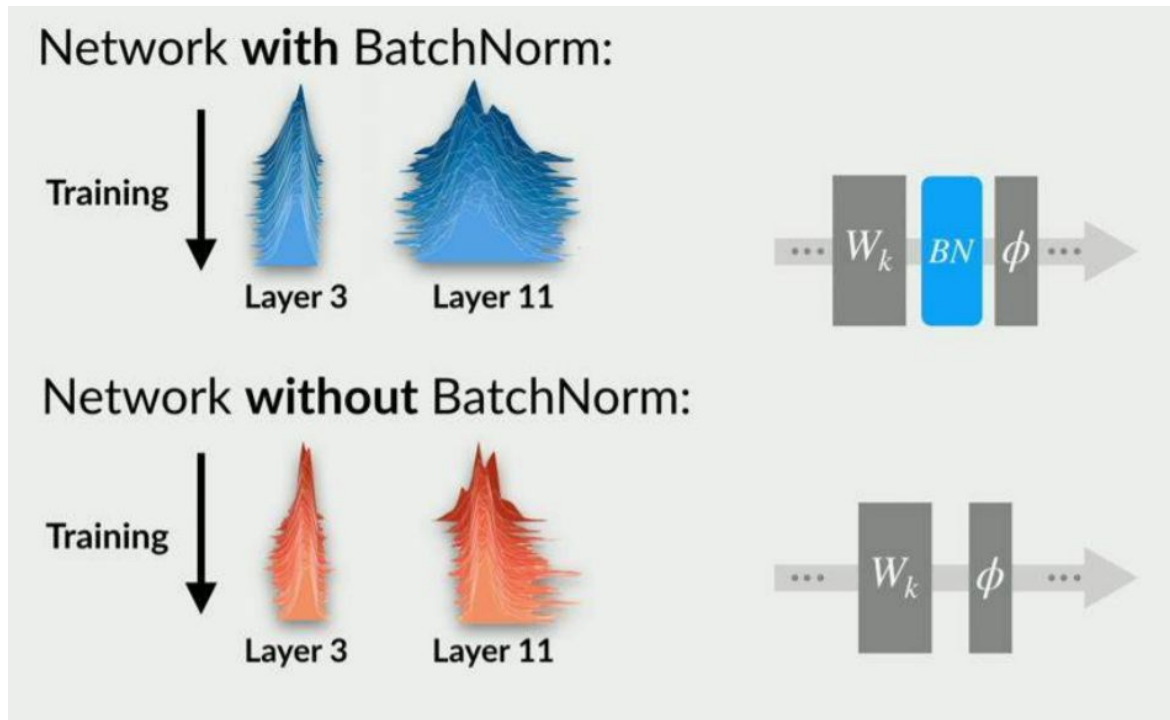





# The internal Covariate Shift Hypothesis



# Experiment: Input Distribution



# Experiment: Noisy Batch Normalization



# Experiment: Noisy Batch Normalization

---

**Algorithm 1** “Noisy” BatchNorm

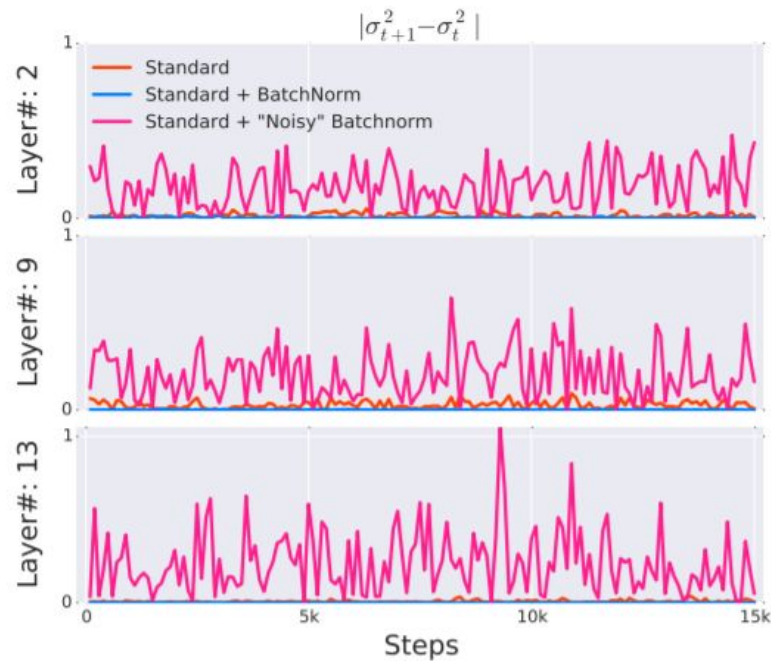

---

```
1: % For constants  $n_m, n_v, r_m, r_v$ .
2:
3: for each layer at time  $t$  do
4:    $a_{i,j}^t \leftarrow$  Batch-normalized activation for unit  $j$  and sample  $i$ 
5:
6:   for each  $j$  do                                      $\triangleright$  Sample the parameters  $(m_j^t, v_j^t)$  of  $D_j^t$  from  $D_j$ 
7:      $\mu^t \sim U(-n_\mu, n_\mu)$ 
8:      $\sigma^t \sim U(1, n_\sigma)$ 
9:
10:    for each  $i$  do                                        $\triangleright$  Sample noise from  $D_j^t$ 
11:      for each  $j$  do
12:         $m_{i,j}^t \sim U(\mu - r_\mu, \mu + r_\mu)$ 
13:         $s_{i,j}^t \sim \mathcal{N}(\sigma, r_\sigma)$ 
14:         $a_{i,j}^t \leftarrow s_{i,j}^t \cdot a_{i,j} + m_{i,j}^t$ 
```


---

In experiments,  $n_\mu = 0.5$ ,  $n_\sigma = 1.25$  and  $r_\mu = r_\sigma = 0.1$


# Experiment: Noisy Batch Normalization




# Experiment: Noisy Batch Normalization




# Experiment: Gradient and Loss Variation



# Experiment: Gradient and Loss Variation



(a) loss landscape




(b) gradient predictiveness


At a particular training step, we measure the variation (shaded region) in loss (a) and L2 changes in the gradient (b) as we move in the gradient direction.



# Theoretical Results



(a) Vanilla Network

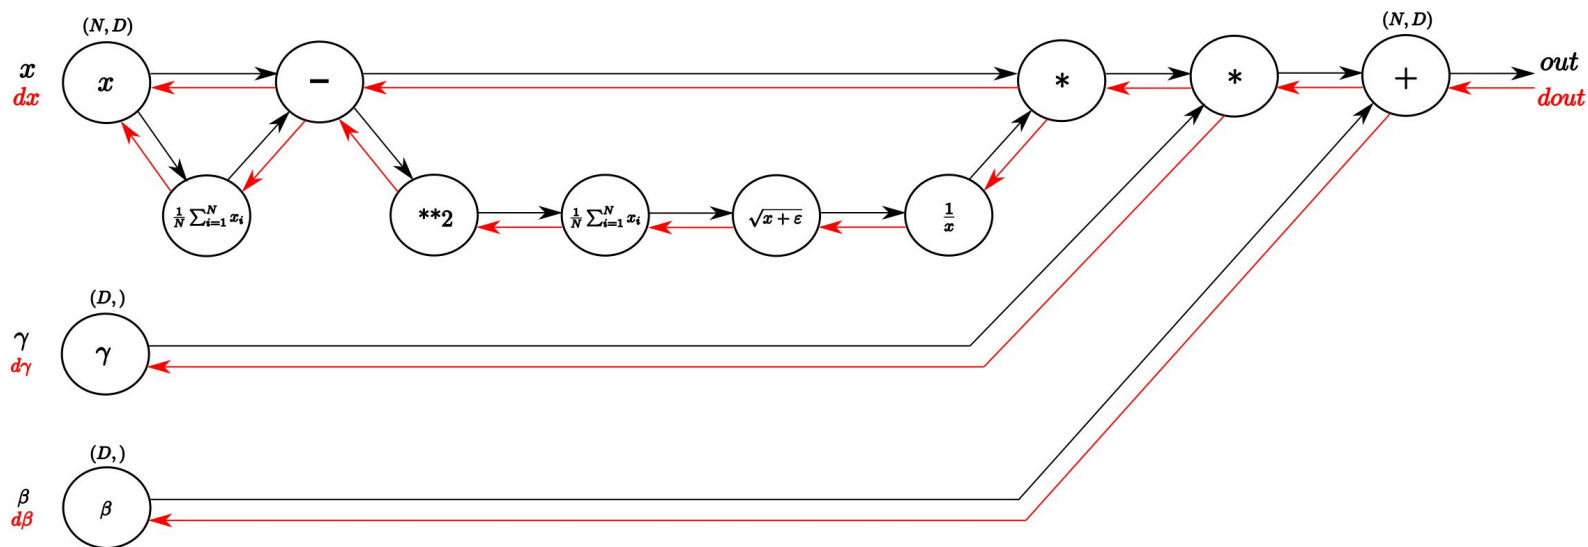


(b) Vanilla Network + BatchNorm Layer


$$|f(x_1) - f(x_2)| \leq L \|x_1 - x_2\|$$

$$\|\nabla_{y_j} \hat{L}\|^2 \leq \frac{\gamma^2}{\sigma_j^2} \left( \|\nabla_{y_j} L\|^2 - \frac{1}{m} \langle 1, \nabla_{y_j} L \rangle^2 - \frac{1}{m} \langle \nabla_{y_j} L, \hat{y}_j \rangle^2 \right)$$

# What is the difference with division by a constant?




# Batch Norm Problems



# Group Normalization

[Wu & He, 2018]



Normalization methods. Each subplot shows a feature map tensor, with  $N$  as the batch axis,  $C$  as the channel axis, and  $(H, W)$  as the spatial axes. The pixels in blue are normalized by the same mean and variance, computed by aggregating the values of these pixels.

# Group Normalization: realization

$$y_i = \gamma \hat{x}_i + \beta$$

$$\hat{x}_i = \frac{1}{\sigma_i} (x_i - \mu_i)$$

$$\mu_i = \frac{1}{m} \sum_{k \in S_i} x_k$$

$$\sigma_i = \sqrt{\frac{1}{m} \sum_{k \in S_i} (x_k - \mu_i)^2} + \epsilon$$

$$S_i = \left\{ k \mid k_n = i_n, \left\lfloor \frac{k_C}{C/G} \right\rfloor = \left\lfloor \frac{i_C}{C/G} \right\rfloor \right\}$$

```
def GroupNorm(x, gamma, beta, G, eps=1e-5):  
    # x: input features with shape [N,C,H,W]  
    # gamma, beta: scale and offset, with shape [1,C,1,1]  
    # G: number of groups for GN  
  
    N, C, H, W = x.shape  
    x = tf.reshape(x, [N, G, C // G, H, W])  
  
    mean, var = tf.nn.moments(x, [2, 3, 4], keep_dims=True)  
    x = (x - mean) / tf.sqrt(var + eps)  
  
    x = tf.reshape(x, [N, C, H, W])  
  
    return x * gamma + beta
```


# Group Normalization: cases

# groups ( $G$ )						
64	32	16	8	4	2	1 (=LN)
24.6	<b>24.1</b>	24.6	24.4	24.6	24.7	25.3
0.5	-	0.5	0.3	0.5	0.6	1.2

# channels per group						
64	32	16	8	4	2	1 (=IN)
24.4	24.5	<b>24.2</b>	24.3	24.8	25.6	28.4
0.2	0.3	-	0.1	0.6	1.4	4.2


Group division. We show ResNet-50's validation error (%) in ImageNet, trained with 32 images/GPU. (Top): a given number of groups. (Bottom): a given number of channels per group. The last rows show the differences with the best number.

# Group Normalization: better than BN?



ImageNet classification error vs. batch sizes. This is a ResNet-50 model trained in the ImageNet training set using 8 workers (GPUs), evaluated in the validation set


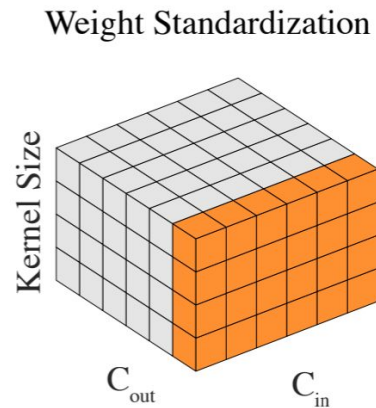
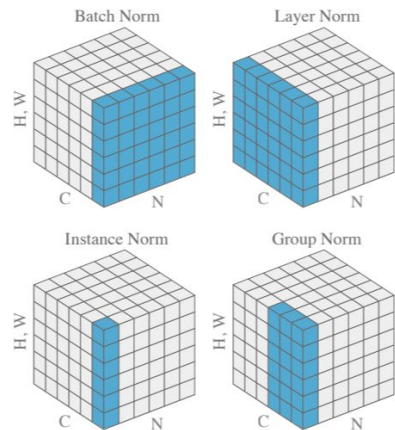
# Group Normalization: better than BN?



Comparison of error curves with a batch size of 32 images/GPU. We show the ImageNet training error (left) and validation error (right) vs. numbers of training epochs. The model is ResNet-50.



# Weight Standardization



[Qiao et al., 2019]

# Weight Standardization

$y = \hat{W} * x$ ;  $\hat{W} \in \mathbb{R}^{O \times I}$  denotes the weights in the layer

$*$  denotes the convolution operation

$$\hat{W} = \left[ \hat{W}_{i,j} \mid \hat{W}_{i,j} = \frac{W_{i,j} - \mu_{W_{i,.}}}{\sigma_{W_{i,.}} + \epsilon} \right]$$

$$\mu_{W_{i,.}} = \frac{1}{I} \sum_{j=1}^I W_{i,j}$$

$$\sigma_{W_{i,.}} = \sqrt{\frac{1}{I} \sum_{j=1}^I (W_{i,j} - \mu_{W_{i,.}})^2}$$

```
class StdConv2d(nn.Conv2d):
```

```
    def forward(self, x):
```

```
        w = self.weight
```

```
        v, m = torch.var_mean(w, dim=[1, 2, 3], keepdim=True, unbiased=False)
```

```
        w = (w - m) / torch.sqrt(v + 1e-10)
```

```
        return F.conv2d(x, w, self.bias, self.stride, self.padding,  
                        self.dilation, self.groups)
```

# Weight Standardization: Lipschitzness


$$\dot{W}_{c,.} = W_{c,.} - \frac{1}{I} 1 \langle 1, W_{c,.} \rangle$$

$$\hat{W}_{c,.} = \dot{W}_{c,.} / \left( \sqrt{\frac{1}{I} \langle 1, \dot{W}_{c,.}^2 \rangle} \right)$$

$$\left| \left| \nabla_{\dot{W}_{c,.}} L \right| \right|^2 = \frac{1}{\sigma_{W_{c,.}}} \left( \left| \left| \nabla_{\hat{W}_{c,.}} L \right| \right|^2 - \frac{1}{I} \langle \hat{W}_{c,.}, \nabla_{\hat{W}_{c,.}} L \rangle^2 \right)$$

$$\left| \left| \nabla_{W_{c,.}} L \right| \right|^2 = \left| \left| \nabla_{\dot{W}_{c,.}} L \right| \right|^2 - \frac{1}{I \cdot \sigma_{W_{c,.}}^2} \langle 1, \nabla_{\hat{W}_{c,.}} L \rangle^2$$

# Weight Standardization: better than BN?



Comparing BN, GN, and WS used with GN on ImageNet and COCO. On ImageNet, BN is trained with large batch sizes while GN and GN+WS are trained with 1 image/GPU. On COCO, BN is frozen for micro-batch training. GN+WS still outperforms both BN and GN comfortably.

	Plain Conv	Weight Std.
Batch Norm.	75.6	75.8
Group Norm.	70.2	<b>76.0</b>

Top-1 accuracy of ResNet-50 trained from scratch on ILSVRC-2012 (BiT-S) with a batch-size of 4096.

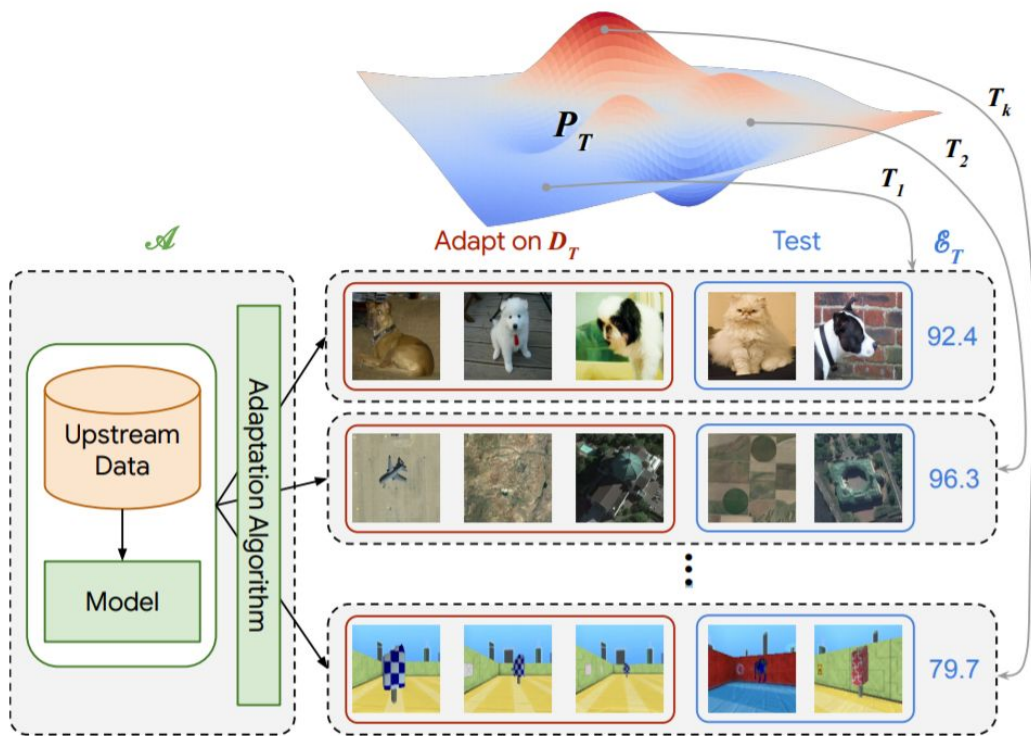
# Downstream Tasks

	BiT-L	Generalist SOTA	Specialist SOTA
ILSVRC-2012	<b>87.54 <math>\pm</math> 0.02</b>	86.4 [57]	88.4 [61]*
CIFAR-10	<b>99.37 <math>\pm</math> 0.06</b>	99.0 [19]	-
CIFAR-100	<b>93.51 <math>\pm</math> 0.08</b>	91.7 [55]	-
Pets	<b>96.62 <math>\pm</math> 0.23</b>	95.9 [19]	97.1 [38]
Flowers	<b>99.63 <math>\pm</math> 0.03</b>	98.8 [55]	97.7 [38]
VTAB (19 tasks)	<b>76.29 <math>\pm</math> 1.70</b>	70.5 [58]	-

	ILSVRC-2012	CIFAR-10	CIFAR-100	Pets	Flowers	VTAB-1k (19 tasks)
BiT-S (ILSVRC-2012)	81.30	97.51	86.21	93.97	89.89	66.87
BiT-M (ImageNet-21k)	85.39	98.91	92.17	94.46	99.30	70.64
Improvement	+4.09	+1.40	+5.96	+0.49	+9.41	+3.77

# VTAB: The Visual Task Adaptation Benchmark

[Zhai et al., 2020]



# VTAB: The Visual Task Adaptation Benchmark

object identification, scene  
classification, pathology detection,  
counting, localization, and 3D  
geometry -> classification

Category	Dataset	Train size	Classes	Reference
● Natural	Caltech101	3,060	102	(Li et al., 2006)
● Natural	CIFAR-100	50,000	100	(Krizhevsky, 2009)
● Natural	DTD	3,760	47	(Cimpoi et al., 2014)
● Natural	Flowers102	2,040	102	(Nilsback & Zisserman, 2008)
● Natural	Pets	3,680	37	(Parkhi et al., 2012)
● Natural	Sun397	87,003	397	(Xiao et al., 2010)
● Natural	SVHN	73,257	10	(Netzer et al., 2011)
● Specialized	EuroSAT	21,600	10	(Helber et al., 2019)
● Specialized	Resisc45	25,200	45	(Cheng et al., 2017)
● Specialized	Patch Camelyon	294,912	2	(Veeling et al., 2018)
● Specialized	Retinopathy	46,032	5	(Kaggle & EyePacs, 2015)
● Structured	Clevr/count	70,000	8	(Johnson et al., 2017)
● Structured	Clevr/distance	70,000	6	(Johnson et al., 2017)
● Structured	dSprites/location	663,552	16	(Matthey et al., 2017)
● Structured	dSprites/orientation	663,552	16	(Matthey et al., 2017)
● Structured	SmallNORB/azimuth	36,450	18	(LeCun et al., 2004)
● Structured	SmallNORB/elevation	36,450	9	(LeCun et al., 2004)
● Structured	DMLab	88,178	6	(Beattie et al., 2016)
● Structured	KITTI/distance	5,711	4	(Geiger et al., 2013)

# Fine-tuning: BiTHyperRule

- SGD + momentum 0.9;
- initial lr = 0.003; decay the learning rate by 10 at 30%, 60% and 90% of the training steps
- Batch Size = 512
- Augmentation: random crops + horizontal flips
- MixUp, with  $\alpha = 0.1$ , for medium and large tasks

Dataset size (number of examples)	Schedule length (fine-tuning)	Use MixUp? (mixing parameter $\alpha=0.1$ )
less than 20K	500 steps	No
20k-500k	10K steps	Yes
more than 500k	20K steps	Yes



Input Image area	1 step: resize to	2 step: random crop to
< 96 x 96	160 x 160	128 x 128
> 96 x 96	448 x 448	384 x 384
R152x4 & > 96 x 96	512 x 512	480 x 480



# mixup: beyond empirical risk minimization

[Zhang et al., 2018]

Image



Label

$[1.0, 0.0]$   
cat dog

$[0.0, 1.0]$   
cat dog

$[0.7, 0.3]$   
cat dog

# mixup: empirical risk minimization (ERM)

$$f \in \mathcal{F}; (x, y) \sim P(X, Y)$$

$$P_\delta(x, y) = \frac{1}{n} \sum_{i=1}^n \delta(x = x_i, y = y_i)$$

$$R(f) = \int l(f(x), y) dP(x, y)$$

$$R_\delta(f) = \int l(f(x), y) dP_\delta(x, y) = \frac{1}{n} \sum_{i=1}^n l(f(x_i), y_i)$$

# mixup: empirical vicinal risk minimization

$$P_\nu(\tilde{x}, \tilde{y}) = \frac{1}{n} \sum_{i=1}^n \nu(\tilde{x}, \tilde{y} | x_i, y_i)$$


$$\nu(\tilde{x}, \tilde{y} | x_i, y_i) = N(\tilde{x} - x_i, \sigma^2) \delta(\tilde{y} = y_i)$$

$$R_\nu(f) = \frac{1}{m} \sum_{i=1}^m l(f(\tilde{x}_i), \tilde{y}_i)$$


$$\lambda \sim \text{Beta}(\alpha, \alpha), \text{ for } \alpha \in (0, \infty)$$

$$\mu(\tilde{x}, \tilde{y} | x_i, y_i) = \frac{1}{n} \sum_j E_\lambda[\delta(\tilde{x} = \lambda \cdot x_i + (1 - \lambda) \cdot x_j, \tilde{y} = \lambda \cdot y_i + (1 - \lambda) \cdot y_j)]$$

# What is mixup doing?



Prediction errors in-between training data. Evaluated at  $x = \lambda + x_i(1-\lambda)x_j$ , a prediction is counted as a “miss” if it does not belong to  $\{y_i, y_j\}$ . The model trained with mixup has fewer misses.

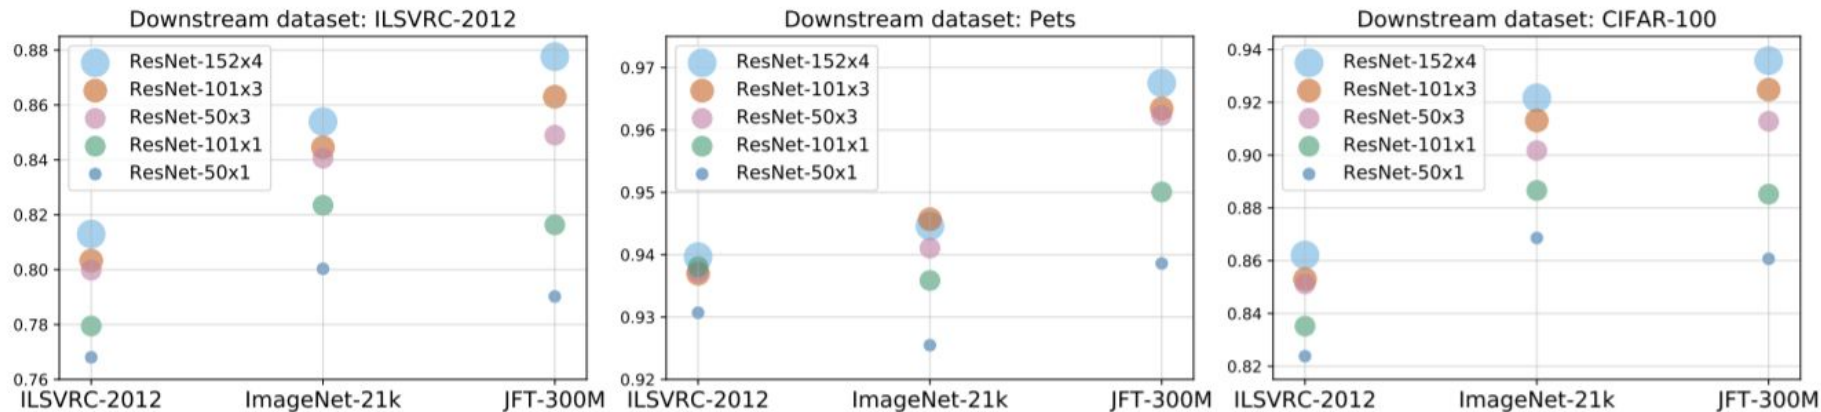


Effect of mixup ( $\alpha = 1$ ) on a toy problem. Green: Class 0. Orange: Class 1. Blue shading indicates  $p(y = 1|x)$ .

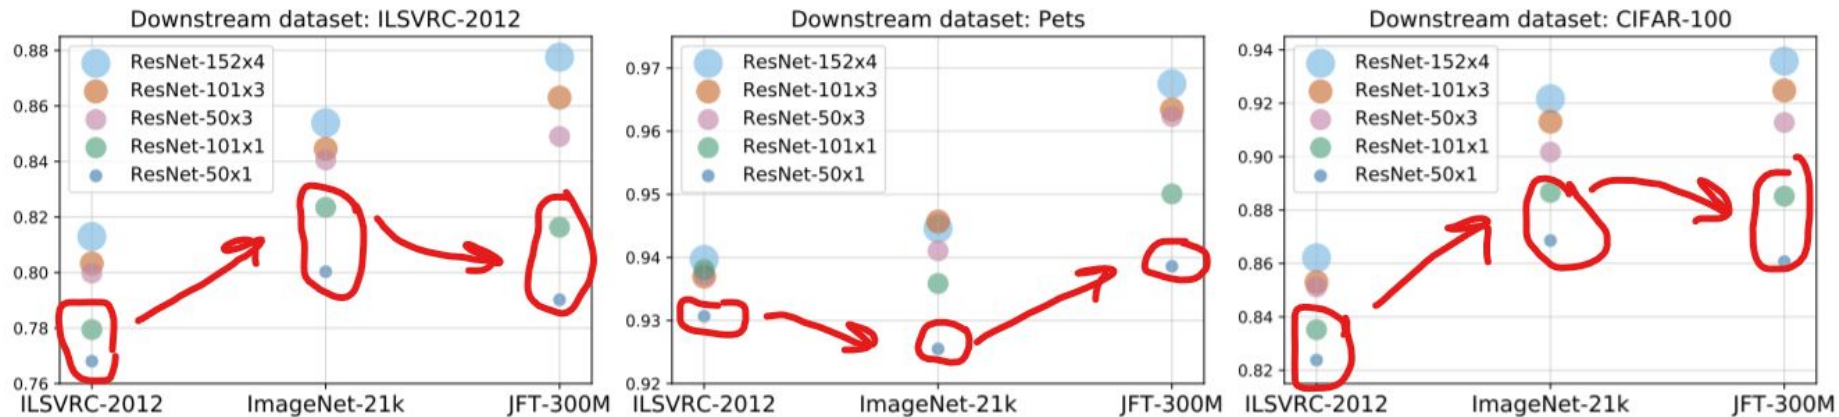
# What is mixup doing?

Model	Method	Epochs	Top-1 Error	Top-5 Error
ResNet-50	ERM (Goyal et al., 2017)	90	23.5	-
	<i>mixup</i> $\alpha = 0.2$	90	<b>23.3</b>	<b>6.6</b>
ResNet-101	ERM (Goyal et al., 2017)	90	22.1	-
	<i>mixup</i> $\alpha = 0.2$	90	<b>21.5</b>	<b>5.6</b>
ResNeXt-101 32*4d	ERM (Xie et al., 2016)	100	21.2	-
	ERM	90	21.2	5.6
	<i>mixup</i> $\alpha = 0.4$	90	<b>20.7</b>	<b>5.3</b>
ResNeXt-101 64*4d	ERM (Xie et al., 2016)	100	20.4	5.3
	<i>mixup</i> $\alpha = 0.4$	90	<b>19.8</b>	<b>4.9</b>
ResNet-50	ERM	200	23.6	7.0
	<i>mixup</i> $\alpha = 0.2$	200	<b>22.1</b>	<b>6.1</b>
ResNet-101	ERM	200	22.0	6.1
	<i>mixup</i> $\alpha = 0.2$	200	<b>20.8</b>	<b>5.4</b>
ResNeXt-101 32*4d	ERM	200	21.3	5.9
	<i>mixup</i> $\alpha = 0.4$	200	<b>20.1</b>	<b>5.0</b>


# BiT Interesting results: Data & Model Sizes



# BiT Interesting results: Data & Model Sizes




# BiT Interesting results: Few-Shot training





# BiT Interesting results: VTAB



# Literature

<https://arxiv.org/abs/1912.11370> - Big Transfer

<https://arxiv.org/pdf/1803.08494.pdf> - Group Normalization

<https://arxiv.org/pdf/1903.10520v1.pdf> - Weight Standardization

<https://arxiv.org/pdf/1710.09412.pdf> - mixup

<https://arxiv.org/pdf/1805.11604.pdf> - How Does Batch Normalization Help Optimization?

[https://www.youtube.com/watch?v=EvAVCxZJN2U&feature=emb\\_logo](https://www.youtube.com/watch?v=EvAVCxZJN2U&feature=emb_logo) - How Does Batch Normalization Help Optimization? (video)